

# J9: Under the Hood of the Next Open Source JVM

Dan Heidinga, J9 VM Interpreter Lead  
Daniel\_Heidinga@ca.ibm.com  
@DanHeidinga  
18 September 2016



# Important disclaimers

- THE INFORMATION CONTAINED IN THIS PRESENTATION IS PROVIDED FOR INFORMATIONAL PURPOSES ONLY.
- WHILST EFFORTS WERE MADE TO VERIFY THE COMPLETENESS AND ACCURACY OF THE INFORMATION CONTAINED IN THIS PRESENTATION, IT IS PROVIDED “AS IS”, WITHOUT WARRANTY OF ANY KIND, EXPRESS OR IMPLIED.
- ALL PERFORMANCE DATA INCLUDED IN THIS PRESENTATION HAVE BEEN GATHERED IN A CONTROLLED ENVIRONMENT. YOUR OWN TEST RESULTS MAY VARY BASED ON HARDWARE, SOFTWARE OR INFRASTRUCTURE DIFFERENCES.
- ALL DATA INCLUDED IN THIS PRESENTATION ARE MEANT TO BE USED ONLY AS A GUIDE.
- IN ADDITION, THE INFORMATION CONTAINED IN THIS PRESENTATION IS BASED ON IBM'S CURRENT PRODUCT PLANS AND STRATEGY, WHICH ARE SUBJECT TO CHANGE BY IBM, WITHOUT NOTICE.
- IBM AND ITS AFFILIATED COMPANIES SHALL NOT BE RESPONSIBLE FOR ANY DAMAGES ARISING OUT OF THE USE OF, OR OTHERWISE RELATED TO, THIS PRESENTATION OR ANY OTHER DOCUMENTATION.
- NOTHING CONTAINED IN THIS PRESENTATION IS INTENDED TO, OR SHALL HAVE THE EFFECT OF:
  - CREATING ANY WARRANT OR REPRESENTATION FROM IBM, ITS AFFILIATED COMPANIES OR ITS OR THEIR SUPPLIERS AND/OR LICENSORS



# Who am I?



- I've been involved with virtual machine development at IBM since 2007 and am now the J9 Virtual Machine Team Lead. J9 is IBM's independent implementation of the JVM.
- I've represented IBM on both the JSR 292 ('invokedynamic') and JSR 335 ('lambda') expert groups and lead J9's implementation of both JSRs.
- I've also maintain the bytecode verifier and deal with various other parts of the runtime.



# J9: IBM's Java Virtual Machine

High performance

High reliability

Serviceability

User application

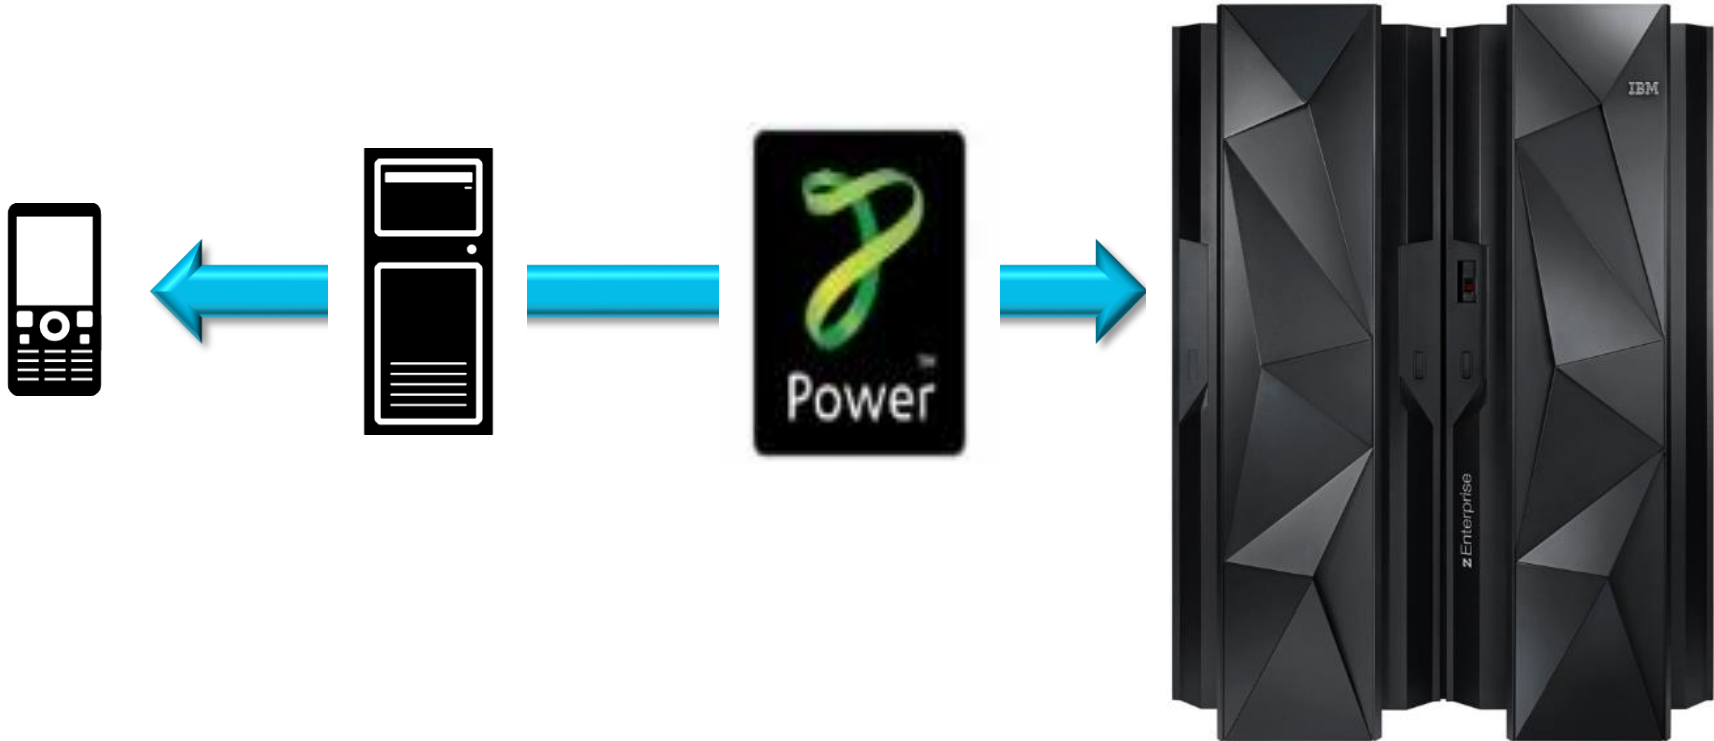
Middle ware

JVM

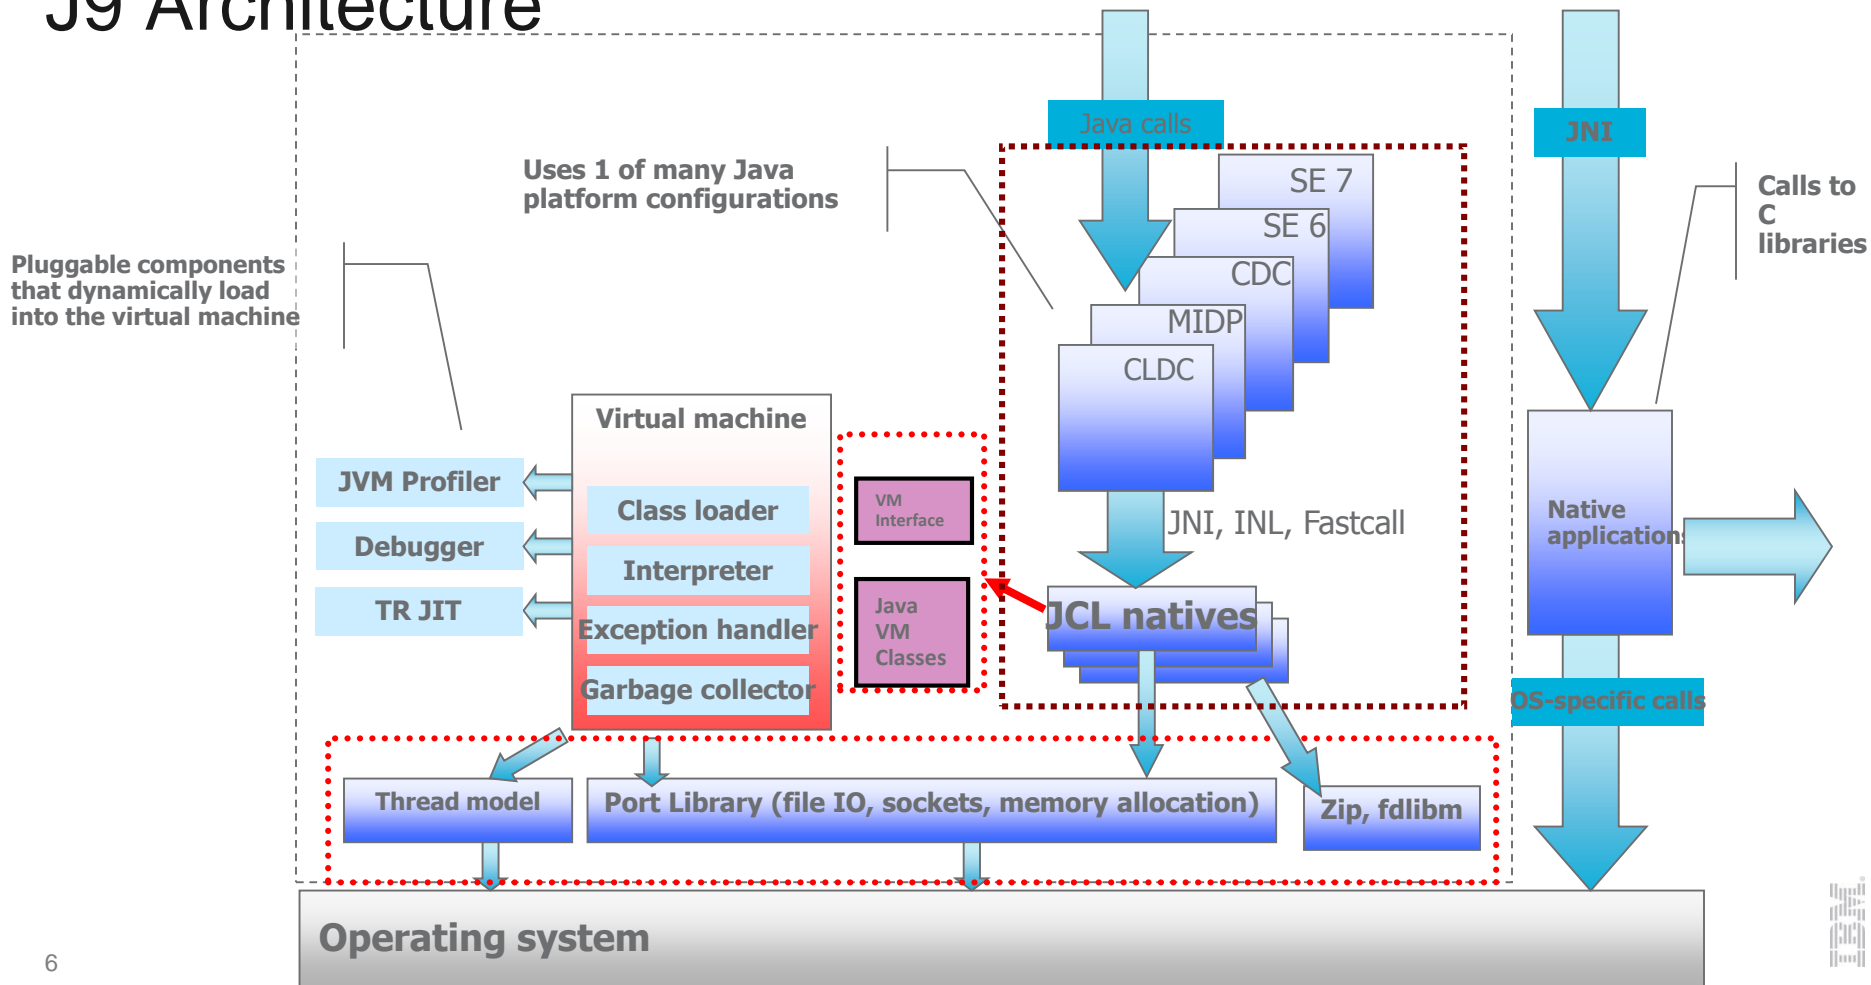
OS



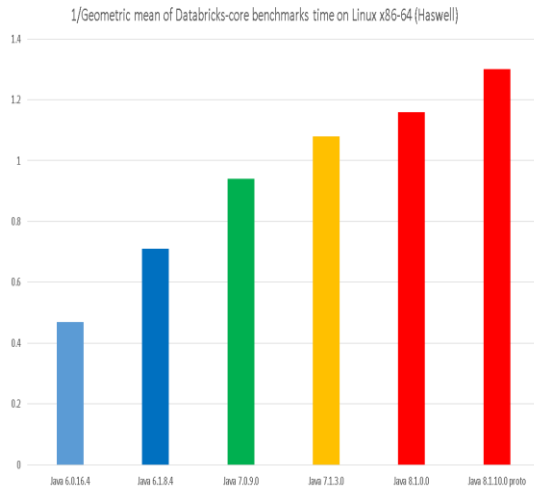
# J9: IBM's Java Virtual Machine scales



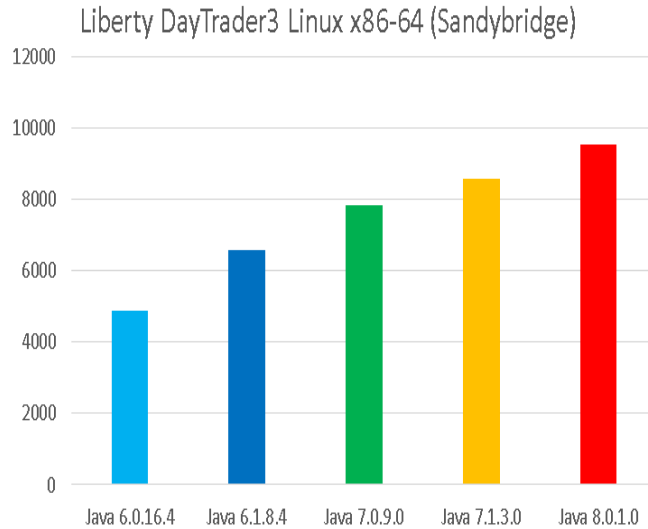
# J9 Architecture



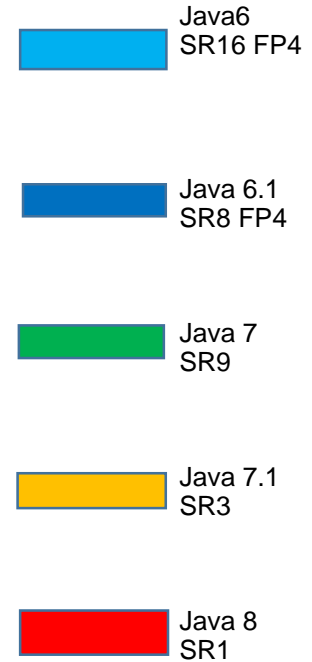
# PERFORMANCE !



Apache Spark 1.4



Daytrader3 Linux Intel



File Edit Classes Applications Categories Methods Info Breakpoints			
<b>Object</b>	<b>Core</b>	(AbtRun-ObsoleteAPI)	addToBeFinalized (public)
AbtAbstractCodePageConve	CLDT	(ACO-API)	<b>finalize (public)</b>
AbtAttributeMapping	CLIM	(CFS-API)	isOSSObject (public)
AbtBase64Coder	CPM	CLDT-API	onFinalizeDo: (public)
AbtCharacterTranslator	AbtNlsKernelApp	(CLIM-API)	removeToBeFinalized (public)
AbtClassElementMapping...	DbgRuntimeFramework	(CT-Internal)	
AbtDependentsCollection	EsAsynchronousCallout	(EM-Internal)	
AbtDOMElementWrapper	K8Types	<b>ES-API</b>	
AbtDOMImplementation	SgmlSupport	ES-Internal	
public	Default: K8VMBuildAPI	instance	all

finalize

"The receiver no longer has any referencers in the image and is being finalized. The default behavior is to do nothing."

^ self

# Envy/Developer: J9's Smalltalk Roots





## Virtual machines on oscilloscopes?

“Lots” of ROM, very little RAM

[https://upload.wikimedia.org/wikipedia/commons/7/7e/Tektronix\\_TDS210\\_Oscilloscope.jpg](https://upload.wikimedia.org/wikipedia/commons/7/7e/Tektronix_TDS210_Oscilloscope.jpg)

Data dominates. If you've chosen the right data structures and organized things well, the algorithms will almost always be self-evident. Data structures, not algorithms, are central to programming.

– Rob Pike



You want me to interpret that?

## 4.1 The `ClassFile` Structure

A class file consists of a single `ClassFile` structure:

```
ClassFile {  
    u4 magic;  
    u2 minor_version;  
    u2 major_version;  
    u2 constant_pool_count;  
    cp_info constant_pool[constant_pool_count-1];  
    u2 access_flags;  
    u2 this_class;  
    u2 super_class;  
    u2 interfaces_count;  
    u2 interfaces[interfaces_count];  
    u2 fields_count;  
    field_info fields[fields_count];  
    u2 methods_count;  
    method_info methods[methods_count];  
    u2 attributes_count;  
    attribute_info attributes[attributes_count];  
}
```

You want me to interpret that?

Constant offsets

Fast access

Easy to find constant\_pool start

Relative offsets

Variable sized arrays

Variable sized entries

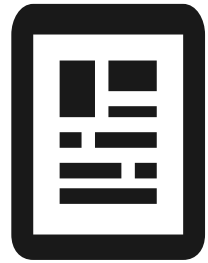
Slow access to specific index

## 4.1 The ClassFile Structure

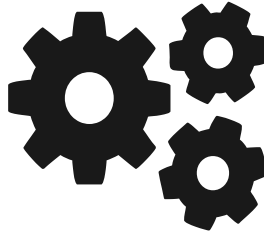
A class file consists of a single ClassFile structure:

```
ClassFile {  
    {  
        u4 magic;  
        u2 minor_version;  
        u2 major_version;  
        u2 constant_pool_count;  
    }  
    cp_info constant_pool[constant_pool_count-1];  
    u2 access_flags;  
    u2 this_class;  
    u2 super_class;  
    u2 interfaces_count;  
    u2 interfaces[interfaces_count];  
    u2 fields_count;  
    field_info fields[fields_count];  
    u2 methods_count;  
    method_info methods[methods_count];  
    u2 attributes_count;  
    attribute_info attributes[attributes_count];  
}
```

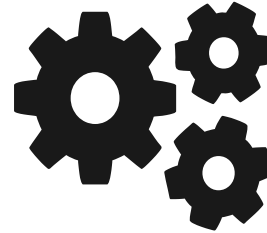
# Conversion to ROM



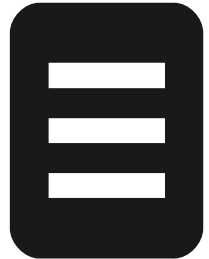
Classfile



Static  
Verification



ROM Class  
Builder



ROMClass



Essential Classfile data

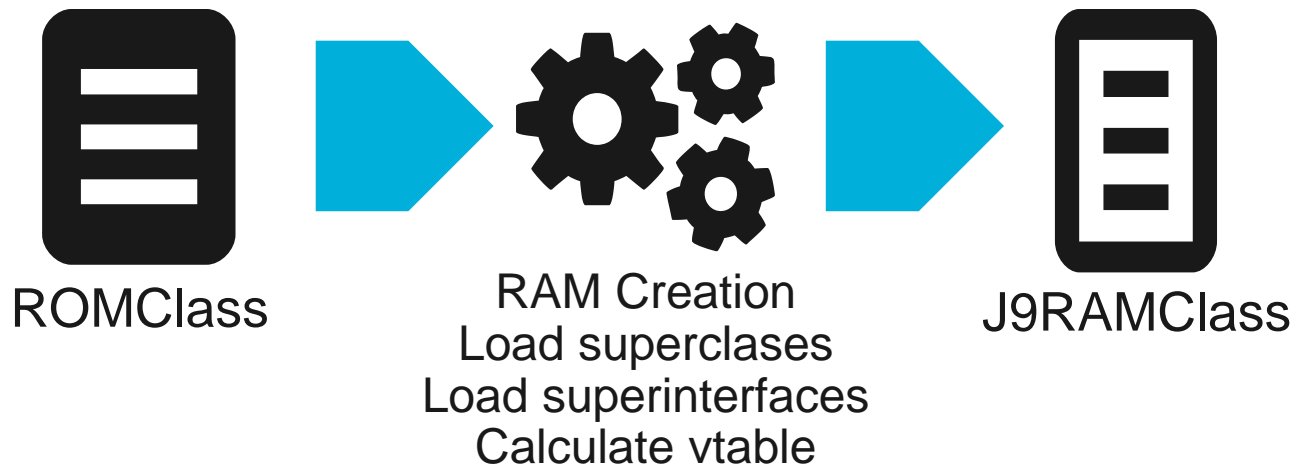
Position independent (SRP)

Less variable than a Classfile

Still variable, but structured

```
J9ROMClass at 0x2b768f835030 {  
  Fields for J9ROMClass:  
    0x0: U32 romSize = 0x0000006B8 (1720)  
    0x4: U32 singleScalarStaticCount = 0x00000000 (0)  
    0x8: J9SRP(struct J9UTF8) className = !j9utf8 0x00002B768F83554E  
    0xc: J9SRP(struct J9UTF8) superclassName = !j9utf8 0x0000000000000000  
    0x10: U32 modifiers = 0x000000021 (33)  
    0x14: U32 extraModifiers = 0x08A00000 (144703488)  
    0x18: U32 interfaceCount = 0x00000000 (0)  
    0x1c: J9SRP(J9SRP(struct J9UTF8)) interfaces = !j9utf8 0x00002B768F8355D6  
    0x20: U32 romMethodCount = 0x00000000D (13)  
    0x24: J9SRP(struct J9ROMMethod) romMethods = !j9rommethod 0x00002B768F8351A8  
    0x28: U32 romFieldCount = 0x00000000 (0)  
    0x2c: J9SRP(struct J9ROMFieldShape) romFields = !j9romfieldshape 0x00002B768F835150  
    0x30: U32 objectStaticCount = 0x00000000 (0)  
    0x34: U32 doubleScalarStaticCount = 0x00000000 (0)  
    0x38: U32 ramConstantPoolCount = 0x00000011 (17)  
    0x3c: U32 romConstantPoolCount = 0x00000011 (17)  
    0x40: J9WSRP(U8) intermediateClassData = 0xB8 (184)  
    0x48: U32 intermediateClassDataLength = 0x0000006B8 (1720)  
    0x4c: U32 instanceShape = 0x00000000E (14)  
    0x50: J9SRP(U32) cpShapeDescription = 0x99991AA0 (2576947872)  
    0x54: J9SRP(struct J9UTF8) outerClassName = !j9utf8 0x0000000000000000  
    0x58: U32 memberAccessFlags = 0x00000000 (0)  
    0x5c: U32 innerClassCount = 0x00000000 (0)  
    0x60: J9SRP(J9SRP(struct J9UTF8)) innerClasses = !j9utf8 0x00002B768F8355D6  
    0x64: U16 majorVersion = 0x0034 (52)  
    0x66: U16 minorVersion = 0x0000 (0)  
    0x68: U32 optionalFlags = 0x000000001 (1)  
    0x6c: J9SRP(U32) optionalInfo = 0x00000013A (314)  
    0x70: U32 maxBranchCount = 0x000000002 (2)  
    0x74: U32 methodTypeCount = 0x00000000 (0)  
    0x78: U32 bsmCount = 0x00000000 (0)  
    0x7c: U32 callSiteCount = 0x00000000 (0)  
    0x80: J9SRP callSiteData = !j9x 0x0000000000000000  
    0x84: U32 classFileSize = 0x00000732 (1842)  
    0x88: U32 classFileCPCCount = 0x000000055 (85)
```

# From ROM to RAM

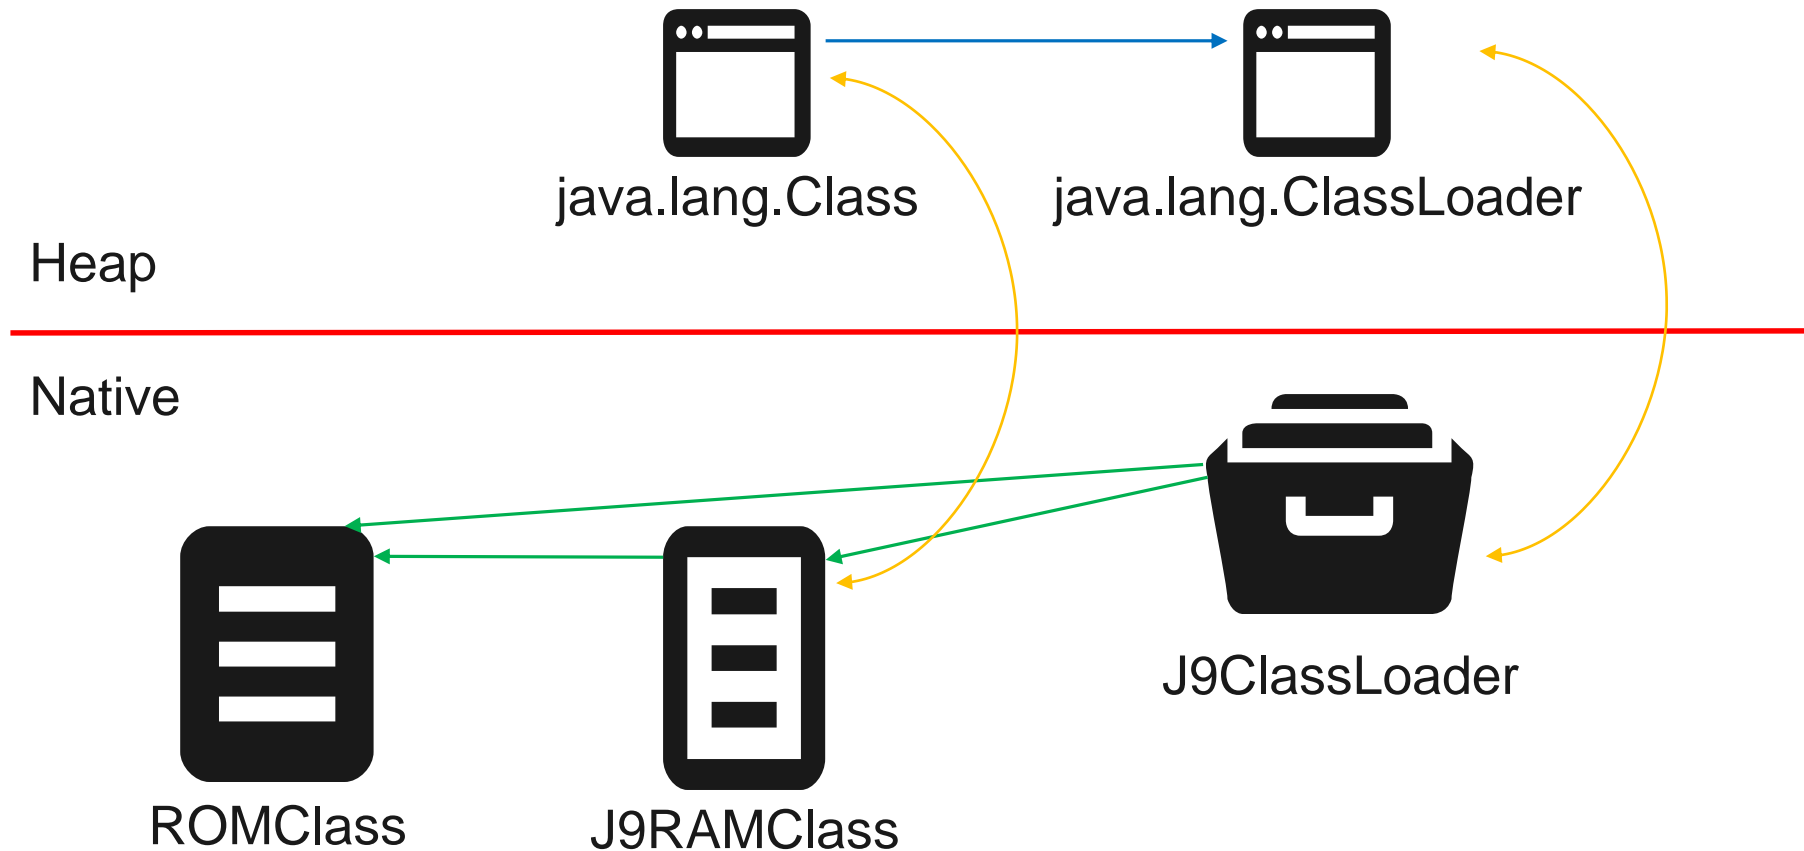


- Live pointers, not SRPs
- Pointer to the ROMClass
- Live data
- Describes the shape of **this** class in **this** VM.
- Requires its entire hierarchy and interfaces to be loaded.

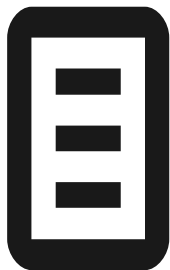
```
j9build@inxem64tvm5:/team/danh/javaone/j9
J9Class at 0x636400 {
  Fields for J9Class:
    0x0: UDATA eyecatcher = 0x0000000099669966 (2573637990)
    0x8: struct J9ROMClass * romClass = !j9romclass 0x00002B768F835030
    0x10: struct J9Class ** superclasses = !j9x 0x0000000000636848
    0x18: UDATA classDepthAndFlags = 0x000000000001E0000 (1966080)
    0x20: U32 classDepthWithFlags = 0x00000000 (0)
    0x24: U32 classFlags = 0x00000000 (0)
    0x28: struct J9ClassLoader * classLoader = !j9classloader 0x00000000167C80B8
    0x30: j9object_t classObject = !j9object 0x00000000E0000200 // java/lang/Class
    0x38: volatile UDATA initializeStatus = 0x0000000000000001 (1)
    0x40: struct J9Method * ramMethods = !j9method 0x0000000000636590 // java/lang/Object.<init>()V
    0x48: UDATA * ramStatics = !j9x 0x0000000000000000
    0x50: struct J9Class * arrayClass = !j9class 0x00000000000064D700 // [Ljava/lang/Object;
    0x58: struct J9Class * packedArrayClass = !j9class 0x0000000000000000
    0x60: UDATA totalInstancesSize = 0x0000000000000004 (4)
    0x68: UDATA packedDataSize = 0x0000000000000000 (0)
    0x70: struct J9ITable * lastITable = !j9itable 0x00002B768A3E4370
    0x78: UDATA * instanceDescription = !j9x 0x0000000000000001
    0x80: UDATA * instanceLeafDescription = !j9x 0x0000000000000001
    0x88: UDATA instanceHotFieldDescription = 0x0000000000000000 (0)
    0x90: struct J9Method * initializerCache = !j9method 0x0000000000000000
    0x98: UDATA romableAotITable = 0x00002B768CAAF64C (47788166149708)
    0xa0: UDATA packageID = 0x00002B768F835031 (47788213882929)
    0xa8: IDATA classPathIndex = 0x0000000000000000 (0)
    0xb0: struct J9Class * subclassTraversalLink = !j9class 0x0000000000751D00 // [Ljava/lang/Thread$State;
    0xb8: struct J9Class * subclassTraversalReverseLink = !j9class 0x0000000000637400 // java/lang/J9VMInternals
    0xc0: void ** itable = !j9x 0x0000000000000000
    0xc8: struct J9PackedFieldsTable * packedFieldsTable = !j9packedfieldstable 0x0000000000000000
    0xd0: UDATA castClassCache = 0x0000000000000000 (0)
    0xd8: void ** jniIDs = !j9x 0x00000000016A72C70
    0xe0: UDATA lockOffset = 0x0000000000000004 (4)
    0xe8: UDATA newInstanceCount = 0x00000000000000BB8 (3000)
    0xf0: IDATA backfillOffset = 0x00000000000000008 (8)
    0xf8: struct J9Class * replacedClass = !j9class 0x0000000000000000
    0x100: UDATA finalizeLinkOffset = 0x0000000000000000 (0)
    0x108: struct J9Class * nextClassInSegment = !j9class 0x0000000000000000
    0x110: UDATA * ramConstantPool = !j9x 0x0000000000636730
    0x118: j9object_t * callSites = !j9x 0x0000000000000000
    0x120: j9object_t * methodTypes = !j9x 0x0000000000000000
    0x128: struct J9Method ** staticSplitMethodTable = !j9x 0x0000000000000000
    0x130: struct J9Method ** specialSplitMethodTable = !j9x 0x0000000000000000
    0x138: struct J9JITExceptionTable * jitMetaDataTable = !j9jitexceptiontable 0x0000000000000000
    0x140: struct J9Class * gcLink = !j9class 0x0000000000000000
}
Class name: java/lang/Object
```



# Classloading: tying it all together



# Bytecode loaded? Check



J9RAMClass



ROMClass

```
stack=2, locals=1, args_size=1
  0: getstatic      #2
  3: ldc             #3
  5: invokevirtual   #4
  8: return
```



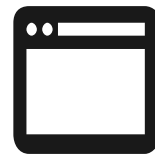
# Threading



Heap  
java.lang.Thread



java.lang.Thread



java.lang.Thread

Native



J9JavaVM



J9VMThread



J9VMThread



J9VMThread



omrthread\_t



omrthread\_t



omrthread\_t

# Per-Thread Interpreter State



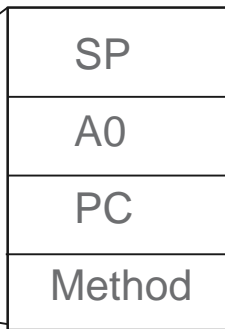
java.lang.Thread



J9VMThread



omrthread\_t



J9Method

ROMMethod

Bytecodes



# Per-Thread Interpreter State



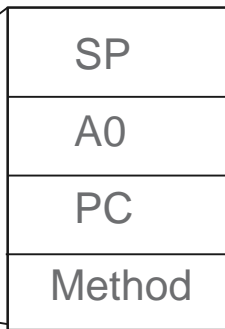
java.lang.Thread



J9VMThread



omrthread\_t



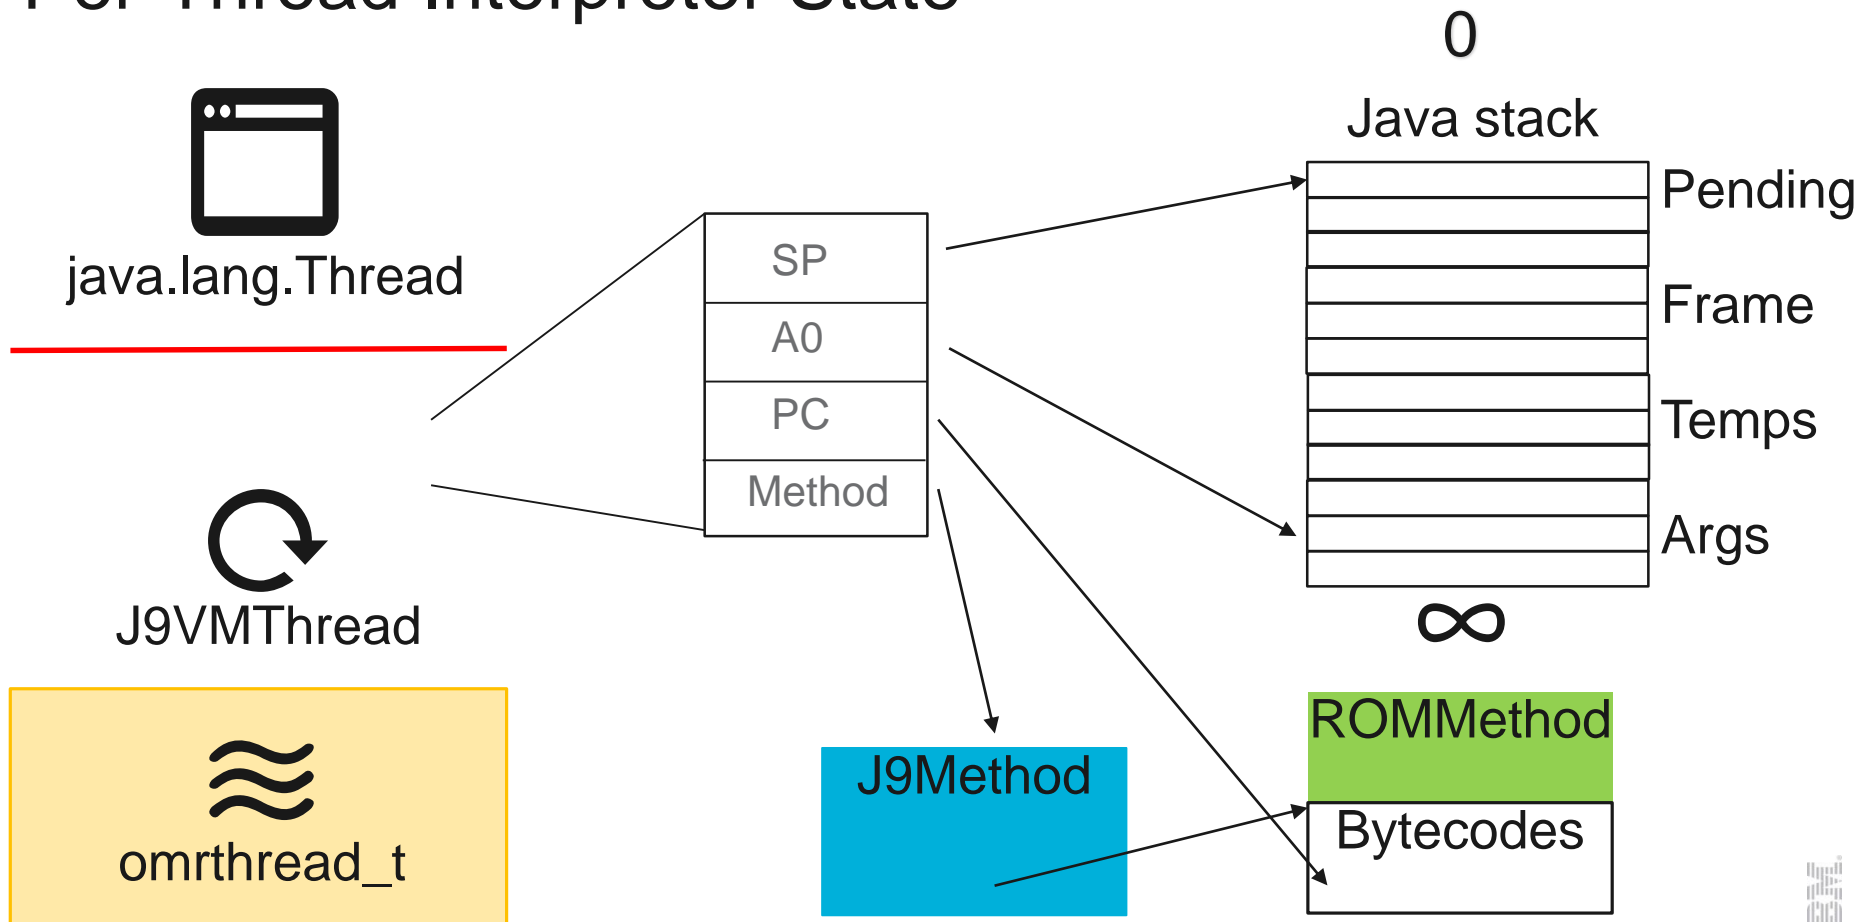
J9Method

ROMMethod

Bytecodes

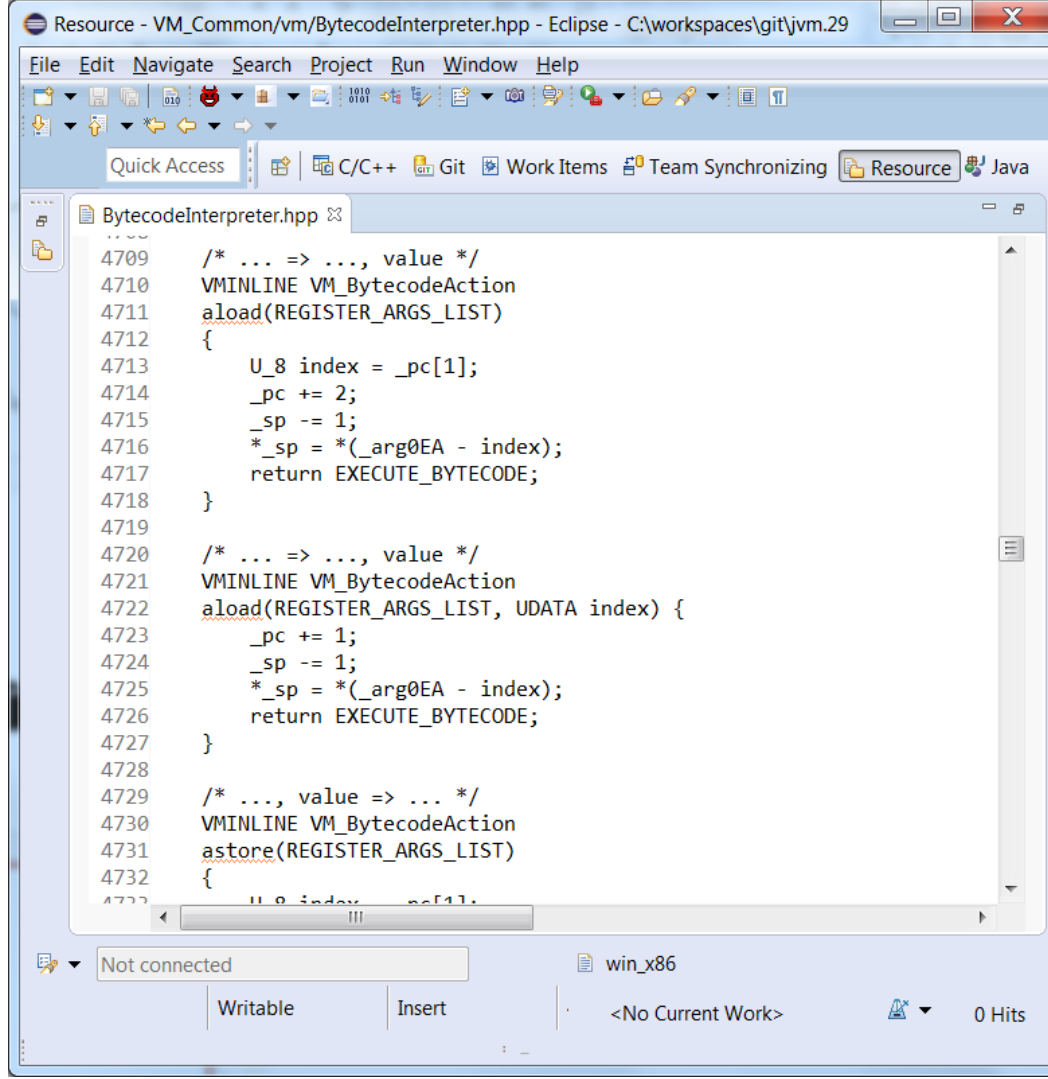


# Per-Thread Interpreter State



# Interpreter

- Written in C++
- Switch statement / computed goto
- Executes:
  - bytecodes
  - INLs
  - builds stack frames
- Transition to the JIT

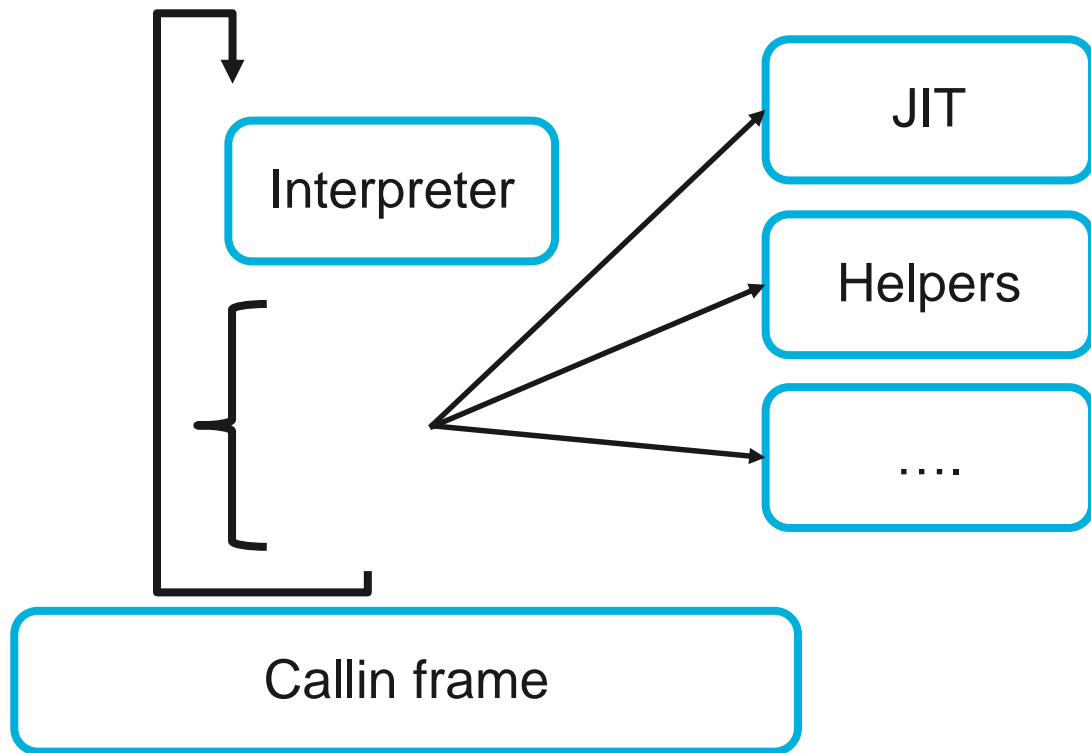


```
Resource - VM_Common/vm/BytecodeInterpreter.hpp - Eclipse - C:\workspaces\git\jvm.29
File Edit Navigate Search Project Run Window Help
Quick Access C/C++ Git Work Items Team Synchronizing Resource Java

BytecodeInterpreter.hpp
4709  /* ... => ..., value */
4710  VM_INLINE VM_BytecodeAction
4711  aload(REGISTER_ARGS_LIST)
4712  {
4713      U_8 index = _pc[1];
4714      _pc += 2;
4715      _sp -= 1;
4716      *_sp = *(_arg0EA - index);
4717      return EXECUTE_BYTECODE;
4718  }
4719
4720  /* ... => ..., value */
4721  VM_INLINE VM_BytecodeAction
4722  aload(REGISTER_ARGS_LIST, UDATA index) {
4723      _pc += 1;
4724      _sp -= 1;
4725      *_sp = *(_arg0EA - index);
4726      return EXECUTE_BYTECODE;
4727  }
4728
4729  /* ..., value => ... */
4730  VM_INLINE VM_BytecodeAction
4731  astore(REGISTER_ARGS_LIST)
4732  {
4733      U_8 index = _pc[1];
```

Not connected | Writable | Insert | win\_x86 | <No Current Work> | 0 Hits

# Primordial Call-in frame





# SharedClasses: ROM pays off

JVM 1



JVM 2



JVM 3

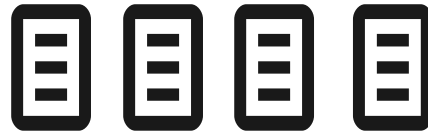


# SharedClasses: ROM pays off

JVM 1



JVM 2



JVM 3



# SharedClasses: ROM pays off

JVM 1



JVM 2



JVM 3



Shared Classes  
Cache



# SharedClasses: “dynamic” AOT

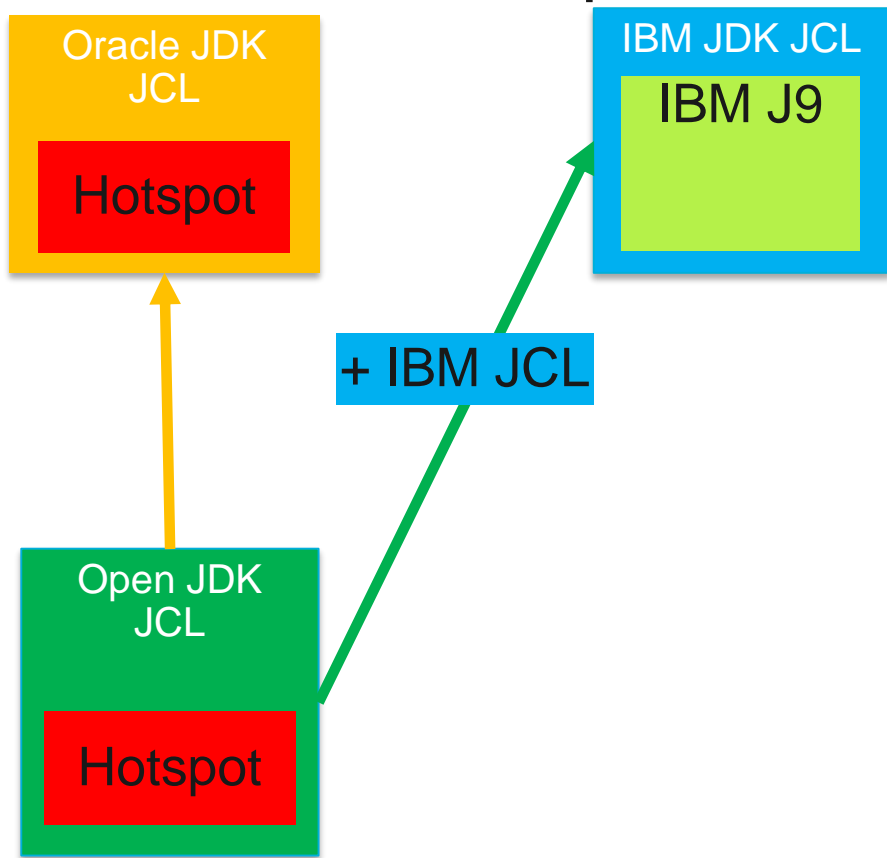


# Diagnostic tools

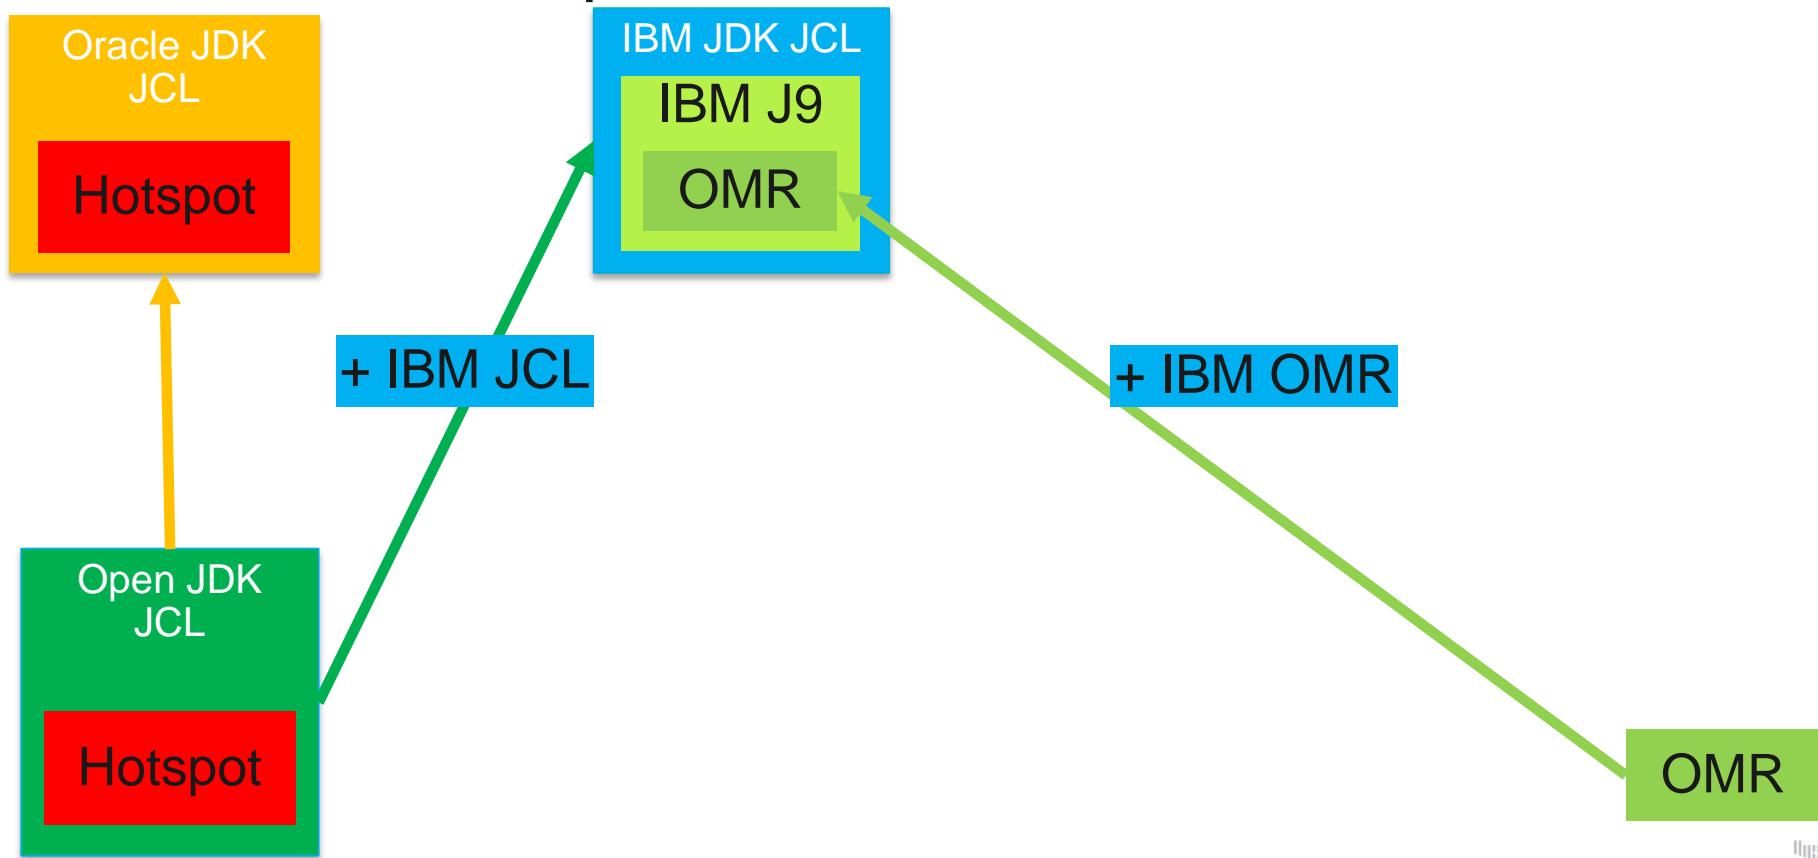
- Diagnostic information (-Xdump)
  - Java dump (javacore)
  - Heap dump
  - System dump (core)
  - JIT dump
- Tracepoints (-Xtrace)
- `bin/jdmpview -core <corefile>`



# J9: Meet IBM's OpenJDK distro



# J9: Meet IBM's OpenJDK & OMR distro



# Eclipse OMR Mission

Build an open reusable language runtime foundation for cloud platforms

- To accelerate advancement and innovation
- In full cooperation with existing language communities
- Engaging a diverse community of people interested in language runtimes
  - Professional developers
  - Researchers
  - Students
  - Hobbyists







# Eclipse OMR

## Created March 2016

<http://www.eclipse.org/omr>  
<https://github.com/eclipse/omr>  
<https://developer.ibm.com/open/omr/>

Dual License:  
Eclipse Public License V1.0  
Apache 2.0

Users and contributors very welcome  
<https://github.com/eclipse/omr/blob/master/CONTRIBUTING.md>



# OMR components

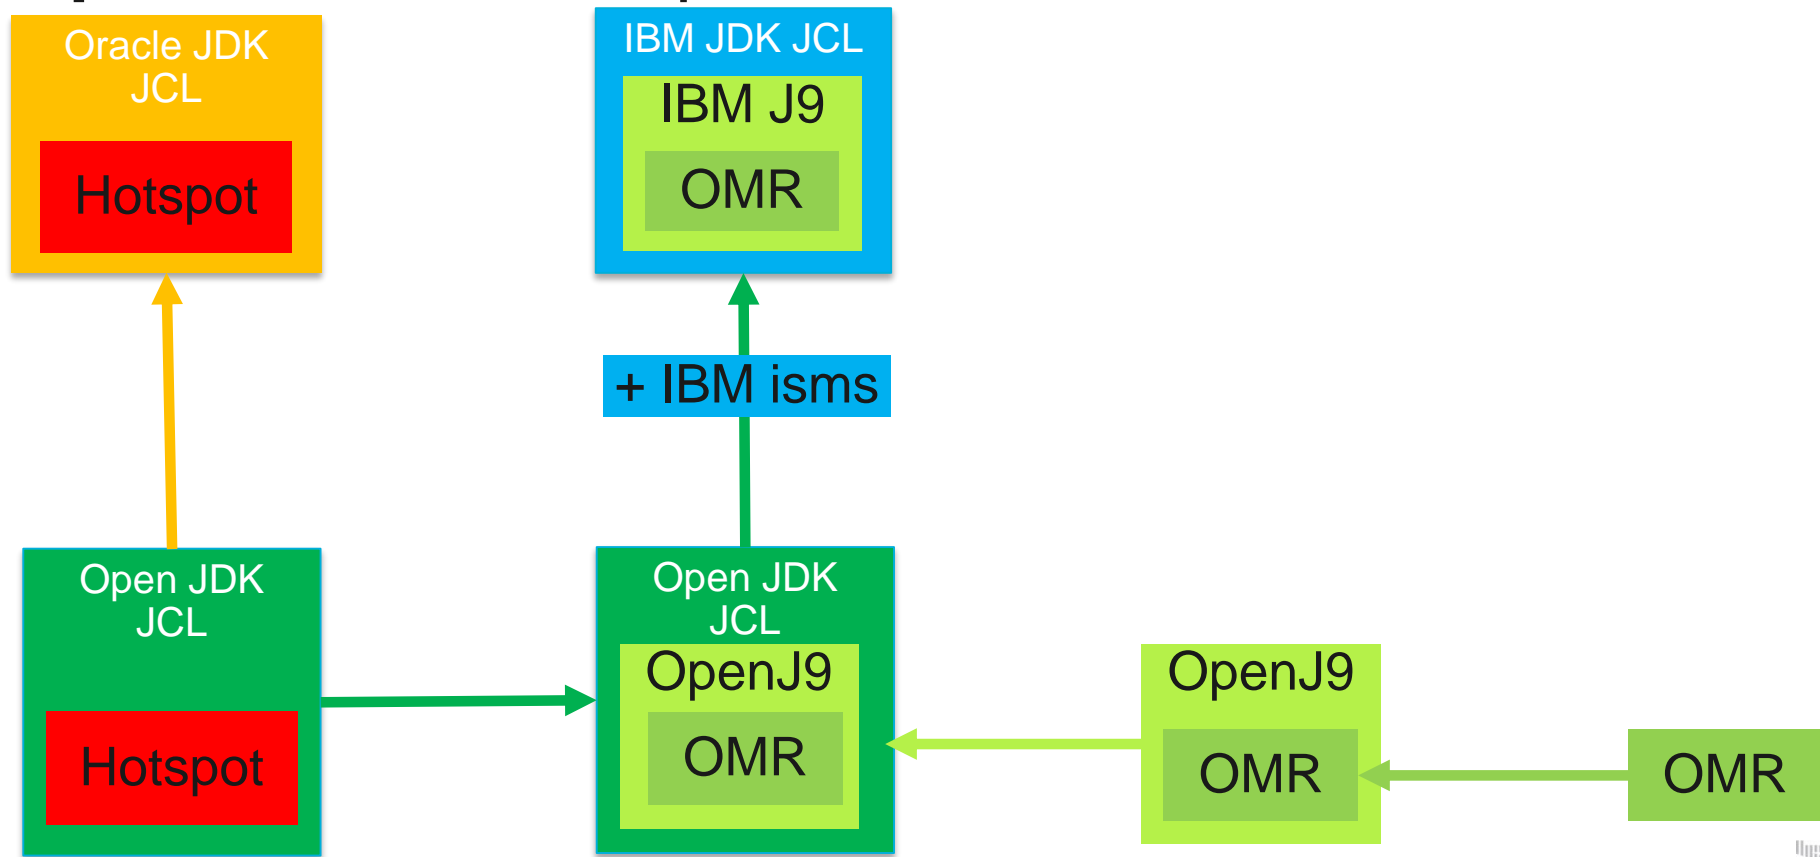
port	platform abstraction (porting) library
thread	cross platform pthread-like threading library
vm	APIs to manage per-interpreter and per-thread contexts
gc	garbage collection framework for managed heaps
compiler	extensible compiler framework
jitbuilder	WIP project to simplify bring up for a new JIT compiler
omrtrace	library for publishing trace events for monitoring/diagnostics
omrsigcompat	signal handling compatibility library
example	demonstration code to show how a language runtime might consume OMR components, also used for testing
fvtest	language independent test framework built on the example glue so that components can be tested outside of a language runtime, uses Google Test 1.7 framework

+ a few others

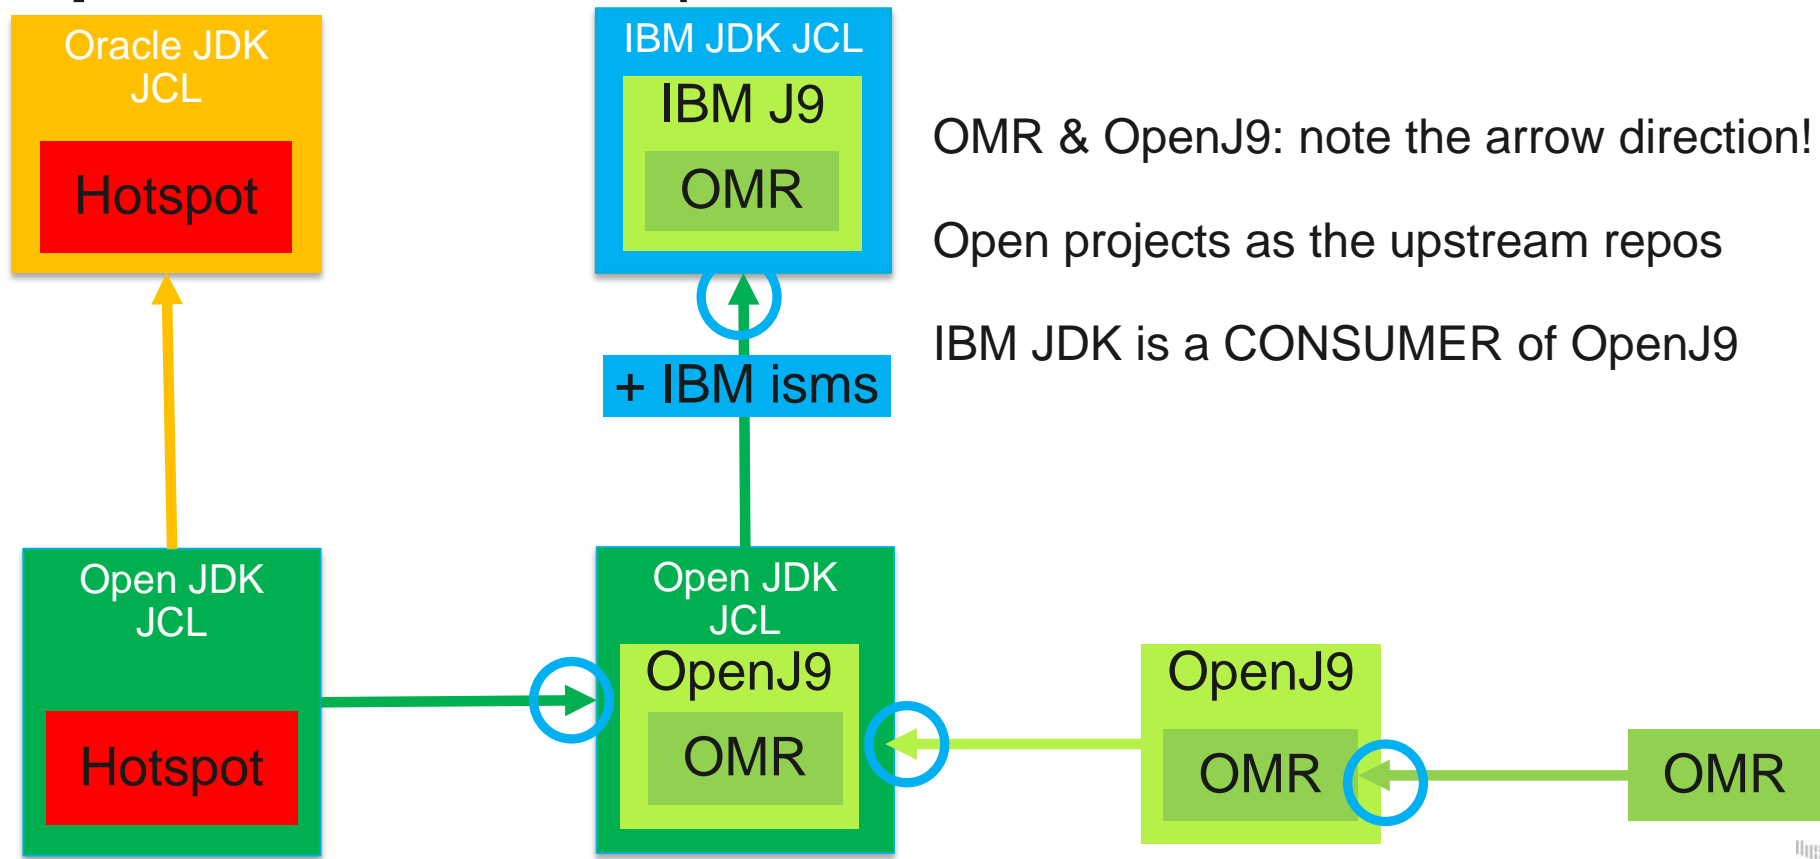
~800KLOC at this point, more components coming!



# OpenJ9: Meet an OpenJDK & OMR distro



# OpenJ9: Meet an OpenJDK & OMR distro



“We believe open ecosystems  
and partnerships are key to our  
future innovation.”

-- Ginni Rometty

(<http://www.ibm.com/annualreport/2014/chairmans-letter.html>)



# Really? Why are you doing this?

- ✓ Collaboration
- ✓ Competition
- ✓ Polyglot
- ✓ Platforms



# When

# Soon



When

Goal: release concurrently  
with Java 9





# Legal Notice

IBM and the IBM logo are trademarks or registered trademarks of IBM Corporation, in the United States, other countries or both.

Java and all Java-based marks, among others, are trademarks or registered trademarks of Oracle in the United States, other countries or both.

Other company, product and service names may be trademarks or service marks of others.

THE INFORMATION DISCUSSED IN THIS PRESENTATION IS PROVIDED FOR INFORMATIONAL PURPOSES ONLY. WHILE EFFORTS WERE MADE TO VERIFY THE COMPLETENESS AND ACCURACY OF THE INFORMATION, IT IS PROVIDED "AS IS" WITHOUT WARRANTY OF ANY KIND, EXPRESS OR IMPLIED, AND IBM SHALL NOT BE RESPONSIBLE FOR ANY DAMAGES ARISING OUT OF THE USE OF, OR OTHERWISE RELATED TO, SUCH INFORMATION. ANY INFORMATION CONCERNING IBM'S PRODUCT PLANS OR STRATEGY IS SUBJECT TO CHANGE BY IBM WITHOUT NOTICE.



# Open source

- IBM is open sourcing its J9 JVM technology
  - Includes Testarossa Just in Time (JIT) compiler
- Eclipse OMR project is leading edge: [github.com/eclipse/omr](https://github.com/eclipse/omr)
  - Project created March 7, 2016 : ~300KLOC
  - NEW! – Compiler contributed September 16, 2016 : ~500KLOC
- Open J9 project is also coming
  - We're working on it at same time as Java 9 development